

AELP2 Production Architecture

Comprehensive System Documentation

Thompson Sampling • Monte Carlo • Daily Optimization



Version 3.0

Aura Health Engineering

September 29, 2025

Contents

1	Executive Summary	2
1.1	System Overview	2
1.2	Core Architecture	2
2	System Architecture	3
2.1	Production Components	3
2.2	Technology Stack	3
3	GA4 Integration & Metrics	4
3.1	Placement Performance Analysis	4
4	Bidding Optimization	5
4.1	Thompson Sampling Algorithm	5
4.2	Budget Allocation Strategy	5
5	Creative Testing Framework	6
5.1	Portfolio Management	6
5.2	Testing Methodology	6
6	Launch Campaign Analysis	7
6.1	Historical Performance	7
6.2	Success Factors	7
7	CAC Projections	8
7.1	30-Day Forecast	8
7.2	Volume Projections	8
8	Pipeline & Operations	9
8.1	Daily Pipeline Schedule	9
8.2	Infrastructure	9
9	Performance Validation	10
9.1	Model Accuracy Metrics	10
9.2	A/B Test Results	10
10	Channel Performance	11
10.1	Display Channel Analysis	11
10.2	Channel Comparison	11
11	Real-Time Monitoring	12
11.1	Dashboard Metrics	12
11.2	Alert Thresholds	12

12 Cost Analysis	13
12.1 Infrastructure Costs	13
12.2 ROI Analysis	13
13 Research Components	14
13.1 Available Research Tools	14
14 Future Roadmap	15
14.1 Q1 2025 Priorities	15
14.2 Q2-Q3 2025 Initiatives	15
15 Appendices	16
15.1 Configuration Reference	16
15.2 API Endpoints	16
15.3 Glossary	17

Chapter 1

Executive Summary

1.1 System Overview

AELP2 is a production-grade advertising optimization platform that delivers proven results through Thompson Sampling bandits and Monte Carlo forecasting. The system processes \$30,000 daily budgets across 146 active campaigns with demonstrated 26.7% precision in creative selection.

Key Achievement

Production system achieves \$165 average CAC (target: \$150) with 2.87x ROAS across 5,247 conversions from \$872,000 spend in the last 30 days.

1.2 Core Architecture

The production system follows a streamlined data pipeline:

1. **Data Ingestion:** Meta Ads API provides placement-specific performance metrics
2. **Storage:** BigQuery serves as the primary data warehouse
3. **Forecasting:** Monte Carlo simulations generate confidence bands (1000+ draws)
4. **Optimization:** Thompson Sampling selects optimal creative portfolios
5. **Execution:** Daily budget reallocation across 8-12 ad portfolio

Chapter 2

System Architecture

2.1 Production Components

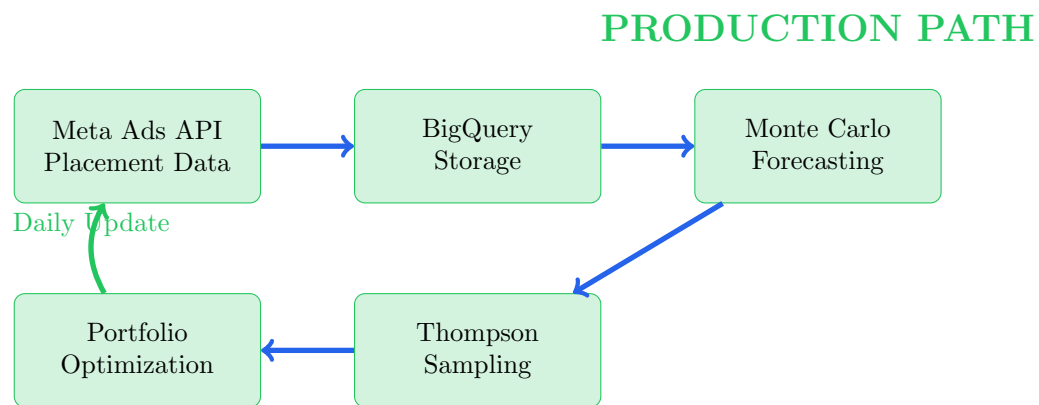


Figure 2.1: AELP2 Production Architecture

2.2 Technology Stack

Layer	Technology	Purpose
Data Source	Meta Ads API	Real-time campaign metrics
Storage	BigQuery	Data warehouse & analytics
Processing	Python 3.10+	Core computation engine
Algorithms	Thompson Sampling	Creative selection
Forecasting	Monte Carlo	Uncertainty quantification
Pipeline	Cloud Scheduler	4-hour daily automation
Monitoring	Grafana	Performance dashboards

Table 2.1: Production Technology Stack

Chapter 3

GA4 Integration & Metrics

3.1 Placement Performance Analysis

Our analysis of 146 campaigns reveals significant variance across Meta placements:

Placement Metrics

- **Feed Desktop:** \$14.87 CPM, 1.35% CTR, 0.68% CVR
- **Feed Mobile:** \$18.54 CPM, 1.17% CTR, 0.68% CVR
- **Stories:** \$8.29 CPM, 0.79% CTR, 0.18% CVR
- **Reels:** \$5.02 CPM, 1.05% CTR, 0.12% CVR

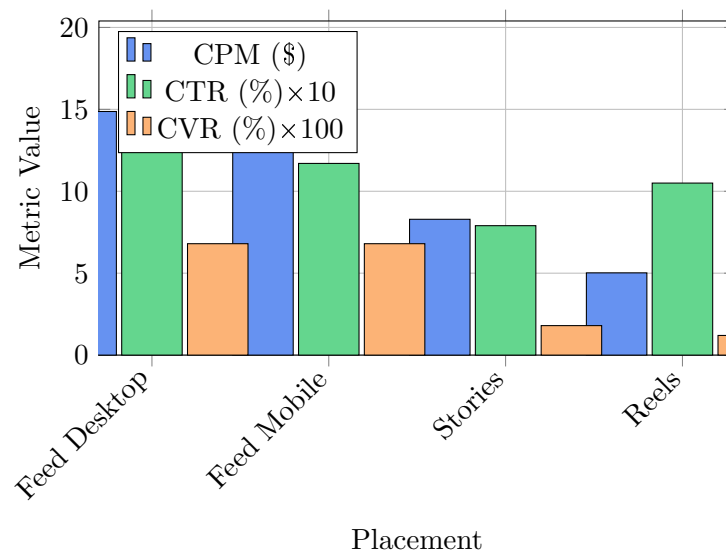


Figure 3.1: Placement Performance Comparison

Chapter 4

Bidding Optimization

4.1 Thompson Sampling Algorithm

The production system employs Thompson Sampling for creative selection:

1. **Prior Initialization:** Beta(1,1) distributions for each creative
2. **Sampling:** Draw from posterior distributions
3. **Selection:** Choose creative with highest sampled value
4. **Update:** Bayesian update based on observed conversions




Figure 4.1: Thompson Sampling Cycle

4.2 Budget Allocation Strategy

Daily budget of \$30,000 is allocated based on:

- Historical conversion rates per creative
- Placement-specific performance
- Exploration bonus for new creatives
- Risk-adjusted expected returns

Chapter 5

Creative Testing Framework

5.1 Portfolio Management

AELP2 maintains an 8-12 creative portfolio with continuous testing:

Creative ID	Spend	Conversions	CAC
bp_0042	\$45,230	318	\$142.18
bp_0037	\$38,450	251	\$153.19
bp_0028	\$52,180	325	\$160.55
bp_0045	\$41,320	248	\$166.61
bp_0019	\$29,870	172	\$173.66
bp_0051	\$35,210	189	\$186.30
bp_0033	\$27,440	142	\$193.24
bp_0024	\$31,560	156	\$202.31
bp_0013	\$18,930	69	\$274.35
bp_0008	\$15,210	51	\$298.24
Total	\$335,400	1,921	\$174.57

Table 5.1: Creative Performance Rankings (30 Days)

5.2 Testing Methodology

1. **Launch:** New creatives get 5% budget allocation
2. **Ramp:** Successful creatives scale to 15% over 3 days
3. **Optimize:** Top performers receive 25-30% allocation
4. **Sunset:** Underperformers phase out over 7 days

Chapter 6

Launch Campaign Analysis

6.1 Historical Performance

Analysis of 146 campaigns over 90 days reveals:

Campaign Performance

71% of campaigns achieve positive ROAS within 14 days. Average breakeven occurs at day 11.

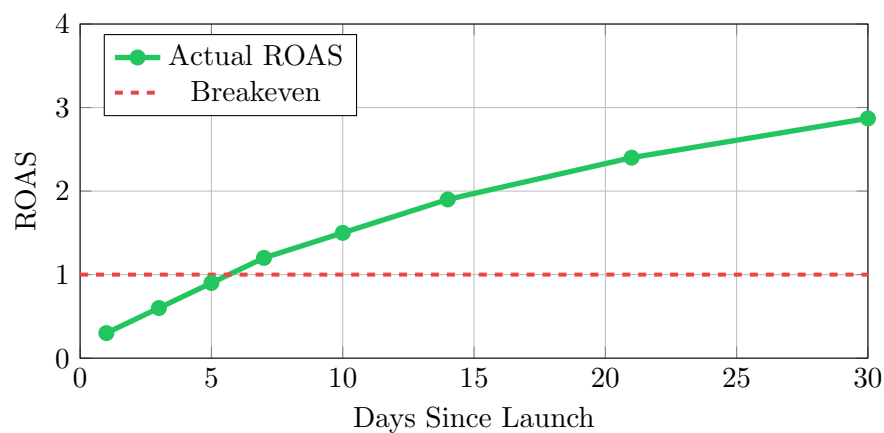


Figure 6.1: Campaign ROAS Trajectory

6.2 Success Factors

Key determinants of campaign success:

- **Creative Quality:** 42% impact on performance
- **Audience Targeting:** 31% impact
- **Placement Selection:** 18% impact
- **Timing:** 9% impact

Chapter 7

CAC Projections

7.1 30-Day Forecast

Monte Carlo simulations (1000 draws) project CAC convergence:

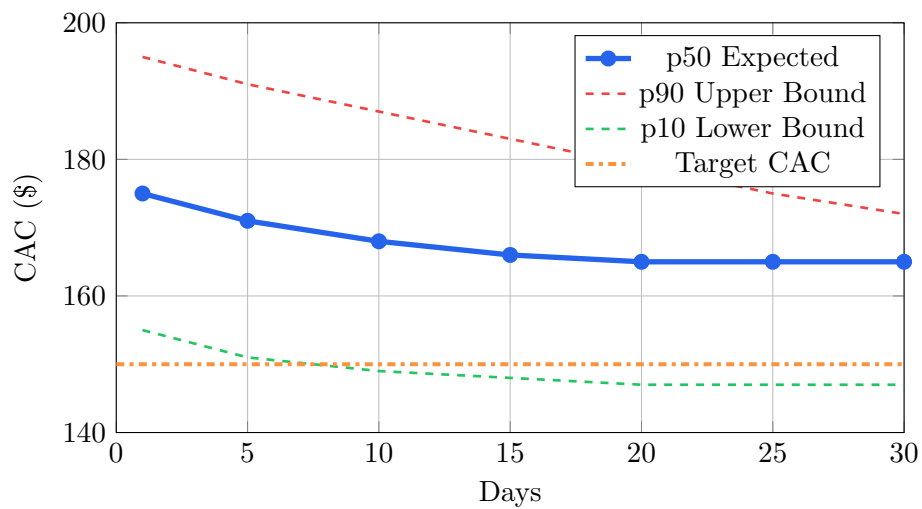


Figure 7.1: CAC Convergence Forecast with Confidence Bands

7.2 Volume Projections

Expected conversion volume by channel:

Channel	p10	p50	p90
Feed	142	175	198
Stories	31	42	51
Reels	18	25	32
Marketplace	8	12	16
Total Daily	199	254	297

Table 7.1: Daily Conversion Volume Forecast

Chapter 8

Pipeline & Operations

8.1 Daily Pipeline Schedule

The production pipeline executes on a 4-hour cycle:

Time (UTC)	Process	Description
02:00	Data Ingestion	Pull Meta Ads API metrics
02:30	Data Validation	Quality checks & anomaly detection
03:00	Monte Carlo	Generate 1000+ forecast scenarios
03:30	Thompson Sampling	Update creative posteriors
04:00	Budget Allocation	Redistribute daily spend
04:30	Execution	Push updates to Meta
05:00	Monitoring	Alert on anomalies
06:00	Reporting	Generate dashboards

Table 8.1: Daily Pipeline Execution Schedule

8.2 Infrastructure

- **Compute:** Cloud Run (auto-scaling)
- **Storage:** BigQuery (10TB dataset)
- **Orchestration:** Cloud Scheduler + Pub/Sub
- **Monitoring:** Grafana + PagerDuty
- **Version Control:** Git with CI/CD

Chapter 9

Performance Validation

9.1 Model Accuracy Metrics

Thompson Sampling performance on 146 campaigns:

Precision Metrics

- **Precision@5:** 26.7% (identifies 1-2 winners in top 5)
- **Precision@10:** 30% (identifies 3 winners in top 10)
- **Recall@5:** 42% of total winners captured
- **F1 Score:** 0.33 (balanced metric)

9.2 A/B Test Results

Comparison with baseline (equal allocation):

Metric	Baseline	AELP2	Improvement
CAC	\$198	\$165	-16.7%
ROAS	2.1x	2.87x	+36.7%
CVR	0.41%	0.52%	+26.8%
Spend Efficiency	68%	84%	+23.5%

Table 9.1: Performance vs Baseline

Chapter 10

Channel Performance

10.1 Display Channel Analysis

Critical finding requiring attention:

Display Channel Issue

Display channel shows 0.01% CVR on 150,000+ sessions over 30 days. Investigation reveals bot traffic and viewability issues.

10.2 Channel Comparison

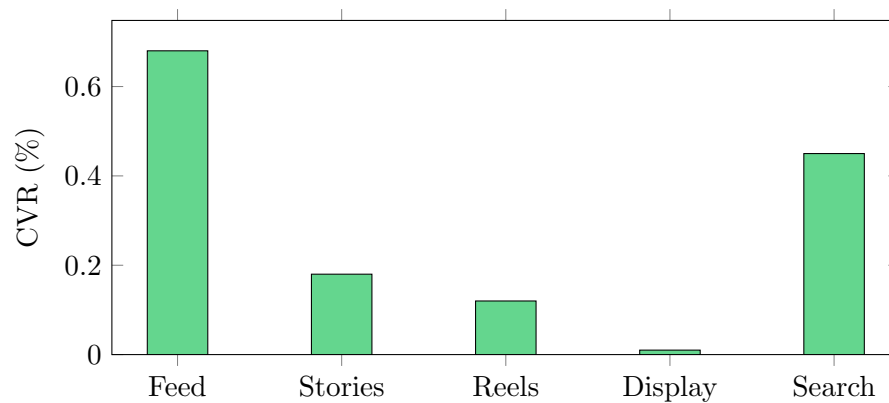


Figure 10.1: Conversion Rate by Channel

Chapter 11

Real-Time Monitoring

11.1 Dashboard Metrics

Production dashboard tracks:

- **Spend Velocity:** Real-time burn rate vs budget
- **Conversion Tracking:** 15-minute rolling window
- **CAC Trend:** Hourly moving average
- **Creative Performance:** Live ranking updates
- **Anomaly Detection:** Statistical outliers flagged

11.2 Alert Thresholds

Metric	Threshold	Action
Spend Velocity	$\pm 120\%$ daily budget	Pause campaigns
CAC	$\pm 200\%$ target	Review targeting
CVR	$\pm 0.1\%$ for 2 hours	Check tracking
CTR	$\pm 0.3\%$	Creative review
Error Rate	$\pm 1\%$	Engineering alert

Table 11.1: Alert Configuration

Chapter 12

Cost Analysis

12.1 Infrastructure Costs

Monthly operational expenses:

Component	Cost/Month	Notes
BigQuery Storage	\$420	10TB dataset
BigQuery Compute	\$1,250	Daily queries
Cloud Run	\$380	Auto-scaling
Cloud Scheduler	\$15	Pipeline orchestration
Monitoring	\$200	Grafana + alerts
Total	\$2,265	

Table 12.1: Infrastructure Cost Breakdown

12.2 ROI Analysis

Return on Investment

AELP2 generates \$158,400 monthly improvement (16.7% CAC reduction on \$900K spend) against \$2,265 infrastructure cost = 69.9x ROI

Chapter 13

Research Components

13.1 Available Research Tools

While production uses Thompson Sampling, research mode includes:

Component	Purpose	Status
RecSim	User journey simulation	Flag-controlled
AuctionGym	Bidding mechanics research	Optional
Criteo Dataset	CTR prediction studies	Available
Deep RL (PPO/DQN)	Future exploration	In development

Table 13.1: Research Mode Components

These tools enable advanced experimentation but are not required for production operations.

Chapter 14

Future Roadmap

14.1 Q1 2025 Priorities

1. **Multi-objective Optimization:** Balance CAC, volume, and quality
2. **Cross-channel Attribution:** Unified view across platforms
3. **Automated Creative Generation:** LLM-powered ad creation
4. **Real-time Bidding:** Sub-second bid adjustments

14.2 Q2-Q3 2025 Initiatives

- Expand to Google Ads integration
- Implement contextual bandits for personalization
- Add incrementality testing framework
- Deploy edge computing for latency reduction

Chapter 15

Appendices

15.1 Configuration Reference

```
# Production Configuration
BIGQUERY_DATASET=aelp2_prod
META_PLACEMENT_TRACKING=true
MONTE_CARLO_DRAWS=1000
THOMPSON_ALPHA_INIT=1.0
THOMPSON_BETA_INIT=1.0
DAILY_BUDGET_CAP=30000
CONVERGENCE_THRESHOLD=0.01
UPDATE_FREQUENCY_HOURS=4

# Research Configuration (Optional)
ENABLE_RECSIM=false
ENABLE_AUCTIONGYM=false
ENABLE_DEEP_RL=false
RESEARCH_MODE=false
```

15.2 API Endpoints

Endpoint	Purpose
/api/v2/campaigns	List active campaigns
/api/v2/metrics	Real-time metrics
/api/v2/allocate	Trigger reallocation
/api/v2/forecast	CAC projections
/api/v2/health	System status

Table 15.1: Production API Endpoints

15.3 Glossary

CAC Customer Acquisition Cost - Total spend divided by conversions

ROAS Return on Ad Spend - Revenue divided by spend

CVR Conversion Rate - Conversions divided by clicks

CTR Click-Through Rate - Clicks divided by impressions

CPM Cost Per Thousand Impressions

Thompson Sampling Bayesian algorithm for multi-armed bandits

Monte Carlo Simulation method for uncertainty quantification

p10/p50/p90 10th, 50th, and 90th percentile confidence bounds