

AELP / AELP2 — System Learnings & First Wave Outputs

Generated 2025-09-29 12:57 UTC

1) Core problem and why

We need to simulate real life response to creatives and placements so the RL agent can learn policies without burning budget. AELP2 turns historical Meta performance and vendor creative evidence into placement aware baselines and calibrated forecasts, then uses an offline bandit simulator to learn allocation policies and output a slate and launch playbook.

2) Architecture — AELP ↔ AELP2 and RL simulator

```
AELP (foundational experiments, legacy pipelines) v AELP2 (production pipelines +
planner + RL sim) | | Meta Insights Vendor/Ad Library v \ /
+-----+ | Ingestion + Normalization |
+-----+ | v +-----+ | Features +
Ranker Scoring | | (new ad p_win, lcb) | +-----+ | v
+-----+ +-----+ | Baselines (by placement) |
---> | Monte Carlo | +-----+ | CAC/Volume | |
+-----+ v | +-----+ v | Offline RL Simulator
| -----> Portfolio policies +-----+ | | Planner UI/API v
| +-----+ v | Launch Playbooks (Meta) | -----> Live
tests & feedback +-----+
```

RL Simulator (offline) ----- Inputs: - Per blueprint priors: p_win, novelty - Baseline draws: CPM, CTR, CVR by placement (triangular p10/p50/p90) - Budget constraints & caps Loop: 1) Sample CTR/CVR/CPM → estimate signups/impressions and CAC 2) Thompson sample success rate per arm (signups/imps proxy) 3) Allocate daily budget to top sampled arms (with caps) 4) Update posteriors using simulated outcomes 5) Emit daily allocations + expected CAC/signups Outputs: - rl_offline_simulation.json (ranking, history) - Portfolio recommendation (8-12 arms + early stop rules)

3) Data pipelines and system flow

- Meta → BQ by placement with backoff/window slicing.
- Vendor/Ad Library via SearchAPI or CSV; normalized into creative_objects.
- Features + new ad ranker scoring (p_win + conformal lcb).
- Baselines (CPM/CTR/CVR) from US paid data; MC forecasts for CAC/volume.
- Offline RL sim (Thompson sampling on signups/imps proxy) to propose a portfolio.
- Planner API/UI packages \$30k/\$50k with playbooks.

4) What we learned (why this design)

- Placement aware baselines reduce drift; CVR clamps avoid outliers.
- Scoring with p_win helps prioritize motifs; conformal lower bounds provide caution.
- Offline RL accelerates exploration safely (policies pre tested); early stop rules protect downside.
- US Ad Library has gaps; augment with SearchAPI/vendor CSV; rely on our own BQ for performance priors.

5) First wave outputs (recommended slate)

Security \$30k/day → signups p50: 153.8, CAC p50 ≈ \$114.9 combined

Balance \$30k/day → signups p50: 368.4

Combined \$60k/day → signups p50: 522.2, CAC p50≈\$114.9, revenue p50≈\$74,972

Security top 8 (id, p_win, su_p50, CAC p50)

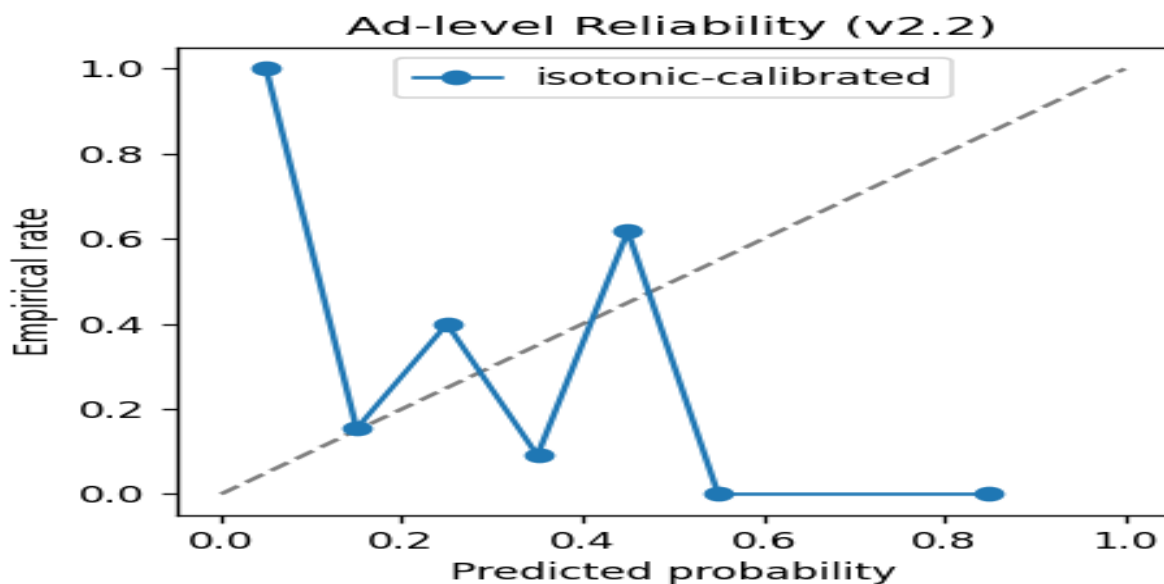
- bp_0042 — p_win 0.222; su_p50 22.7; CAC \$166
- bp_0113 — p_win 0.222; su_p50 22.4; CAC \$167
- bp_0114 — p_win 0.222; su_p50 22.4; CAC \$168
- bp_0117 — p_win 0.222; su_p50 22.6; CAC \$166
- bp_0116 — p_win 0.116; su_p50 18.2; CAC \$206
- bp_0041 — p_win 0.092; su_p50 17.5; CAC \$214
- bp_0001 — p_win 0.000; su_p50 14.0; CAC \$267
- bp_0002 — p_win 0.000; su_p50 14.1; CAC \$266

Balance top 8 (id, p_win, su_p50, CAC p50)

- bpbal_0001 — p_win 0.706; su_p50 45.9; CAC \$82
- bpbal_0002 — p_win 0.706; su_p50 46.1; CAC \$81
- bpbal_0004 — p_win 0.706; su_p50 46.2; CAC \$81
- bpbal_0006 — p_win 0.706; su_p50 46.5; CAC \$81
- bpbal_0009 — p_win 0.706; su_p50 46.3; CAC \$81
- bpbal_0010 — p_win 0.706; su_p50 45.7; CAC \$82
- bpbal_0014 — p_win 0.706; su_p50 45.6; CAC \$82
- bpbal_0017 — p_win 0.706; su_p50 46.1; CAC \$81


6) Accuracy & coverage

Ranker precision@5: 0.2666666666666667, precision@10: 0.3



7) Example creatives (thumbnails)

cool_preview_v1.jpg



orig_freeze_the_chaos_2.jpg



orig_spot_the_scam_1.jpg



8) Next steps

- Complete 90 day by placement backfill and recalibrate.
- Expand Balance set and run offline RL across offer variants.
- Launch controlled live test with early stop rules and measure lift vs. historical CAC.